

AI VALUE BLUEPRINT

Testing Guide for Consultants

How to run a full test of the Blueprint pipeline

using the Meridian Talent Partners test kit

AI Assist BG | Internal Use Only | May 2026

Purpose of this guide: This document tells you — step by step — how to use the Meridian Talent Partners test files to run a full Blueprint pipeline test. You do not need any technical knowledge. Everything you need to type is written out for you. Read this guide from start to finish before you begin.

What is in this guide:

Section	Title	Page
Section 1	Before You Start — What You Have and What It Is	2
Section 2	How to Start a Blueprint Test in Claude	3
Section 3	Step-by-Step: Running All 5 Pipeline Steps	4
Section 4	The Quality Gates — What They Mean and What to Do	9
Section 5	What Good Output Looks Like — Your Checklist	11
Section 6	What to Do When Something Goes Wrong	13
Section 7	Four Test Scenarios to Try	15
Section 8	Recording Your Results	16

Section 1: Before You Start — What You Have and What It Is

You have been given a test kit with 10 files. These files represent a fictional company called **Meridian Talent Partners** — a recruitment firm based in Sofia, Bulgaria. The company and all its data are completely made up, but everything is designed to look and feel like a real client engagement. This allows you to test the full Blueprint pipeline without using any real client data.

The 10 test files

File	What It Is	Required?
00_questionnaire_answers.md	All 30+ intake form answers for Meridian. This is what a real client fills in before you start the Blueprint.	Required
01_financial_summary_FY2025.pdf	Full Profit & Loss statement, revenue breakdown by service line, and financial KPIs for FY2025.	Required
02_org_chart_team_structure.pdf	Company structure: all 6 departments, headcount per role, and leadership team bios.	Required
03_sales_pipeline_data.pdf	Sales pipeline data: stage conversion rates, time-to-fill analysis, top clients, BD activity.	Required
04_process_documentation_sop.pdf	4 core recruitment process SOPs with known issues and adherence rates noted.	Required
05_marketing_customer_data.pdf	Client satisfaction survey results (March 2026), client base analysis, and marketing performance.	Optional
06_technology_inventory.pdf	Full list of every tool and system Meridian uses, with data quality assessment per data domain.	Optional
07_strategic_plan_FY2026.pdf	CEO's confidential strategic plan for FY2026: 4 priorities, financial targets, and key risks.	Optional
08_previous_ai_initiatives.pdf	6 past AI and digital initiatives with honest post-mortems — what was tried, what failed, what was learned.	Optional
README_test_kit.md	Technical reference for developers and administrators. You do not need to read this.	Reference only

FILES The 4 Required files (01–04) must always be uploaded. The 4 Optional files (05–08) make the test richer — include them unless you are specifically testing what happens with thinner materials (see Section 7 for that scenario).

The fictional company at a glance

Detail	Value
Company name	Meridian Talent Partners OOD
Industry	Recruitment & Talent Acquisition (Professional Services)
Location	Sofia, Bulgaria (HQ) + Bucharest, Romania + Warsaw, Poland
Revenue	€4,823,400 (FY2025) — growing 11.6% year-on-year

Detail	Value
Employees	68 full-time staff across 6 departments
Core business	Permanent placement, executive search, contract staffing, RPO pilots
Key problem	Too slow to fill roles (38 days avg vs 28-day target); candidate database unusable; manual admin consuming consultant time
AI readiness	Mixed — good strategy, poor data quality, GDPR gaps, some internal resistance from senior staff
Budget signal	In discussion — €10K–€50K range being considered

Section 2: How to Start a Blueprint Test in Claude

The Blueprint pipeline runs entirely inside Claude (claude.ai). Each of the 5 steps is triggered by uploading the right files and typing a specific instruction. This section explains how to set up your Claude session correctly before you begin.

Step-by-step: Opening Claude and preparing your session

S te p	Action	Detail
1	Open Claude	Go to claude.ai and log in with your AI Assist BG account. Use the standard Claude Sonnet model — do not switch models during a test run.
2	Start a new conversation	Click "New Chat" at the top left. Each Blueprint test run should be in a fresh conversation. Do not continue a test run in an existing chat.
3	Keep all 10 files ready	Have all the test files open in a folder on your computer so you can upload them quickly. You will upload different files at different steps — do not upload everything at once at the start.
4	Open the questionnaire answers	Open the file called 00_questionnaire_answers.md in any text editor (Notepad, Word, or a browser). You will copy-paste the whole document content during Step 1.
5	Note your start time	Write down the time you begin. The full pipeline takes 3–4.5 hours if everything runs smoothly. Most of that time is Claude working — you are waiting and reviewing, not typing continuously.

■ **IMPORTANT**

The Blueprint pipeline uses a project in Claude that has the skill files already loaded. Make sure you are working inside the correct Claude project — not a generic Claude chat. If you are unsure, ask your manager which project to use before starting.

Section 3: Step-by-Step — Running All 5 Pipeline Steps

The Blueprint pipeline has 5 steps. You run them in order. After each step, Claude produces an output that feeds into the next step. Do not skip steps and do not move to the next step until the current one is complete and you have reviewed it. Each step shows exactly what to upload and exactly what to type.

PRE

Before Step 1

Run the pre-pipeline quality screen

Before running Step 1, tell Claude to assess the quality of the materials.

What to upload: Upload all 9 content files at once (files 01–08 plus file 00 the questionnaire answers).

What to type:

"I am running a Blueprint pipeline test. The client is Meridian Talent Partners OOD, engagement ID BP-TEST-001. I have uploaded the client's intake form responses and all 8 supporting documents. Please run the pre-pipeline quality screen: score the intake form, assess material sufficiency, and tell me if we should proceed."

What to expect: Claude confirms all files were received, scores each of the 7 form sections (0–3 points each), calculates a Form Quality Score, and gives a material sufficiency rating (Sufficient / Marginal / Insufficient). For this test kit you should get ~95% (Excellent) and a decision of PROCEED.

1

Intake Analysis

Produces the Compressed Client Dossier (3–4 pages, internal)

What this step does: Claude reads all the client materials and produces a structured internal dossier summarising who the client is, their key data points, pain points in priority order with evidence, and hypotheses for AI opportunities.

What to upload: All files should already be uploaded from the pre-pipeline step. If not, upload all 9 files now.

What to type:

"Please run Step 1 of the Blueprint pipeline — the intake analysis. Produce the full Compressed Client Dossier for Meridian Talent Partners OOD using all the uploaded materials."

What to expect: A 3–4 page dossier with 8 sections (A through H). Every important claim should carry a confidence tag like [Document-Backed] or [Form-Stated]. The dossier ends with a "Quality Self-Assessment" block showing estimated scores across 5 dimensions.

After Claude finishes — trigger the Quality Gate by typing:

"Please run Quality Gate QG-1 on this dossier and tell me the result."

2

Maturity Scoring

Produces the AI Readiness Snapshot (1 page, client-facing)

What this step does: Claude reads the dossier from Step 1 and scores the client across 6 dimensions (Strategy, Data, Technology, People, Processes, Governance). Each dimension gets a level — Early, Developing, or Established — with a written rationale.

What to upload: Nothing new — the dossier is already in the conversation.

What to type:

"Please run Step 2 — the maturity scoring. Use the Compressed Client Dossier you just produced to score Meridian Talent Partners across all 6 AI readiness dimensions. Produce the full AI Readiness Snapshot."

What to expect: A scorecard table showing all 6 dimensions with their levels and key evidence. Then a rationale paragraph per dimension, an "Overall Pattern" paragraph, and "Key Constraints" bullets. Ends with Quality Self-Assessment block.

After Claude finishes — trigger QG-2:

"Please run Quality Gate QG-2 on the AI Readiness Snapshot."

3

Opportunity Harvesting

Produces the Scored Opportunity Map (4–6 pages, client-facing)

What this step does: Claude generates 5–7 specific AI opportunities for Meridian, scores each on Impact, Feasibility, and Alignment (each 1–5), and classifies them as Quick Wins, Foundation Builders, or Big Bets.

What to upload: Nothing new.

What to type:

"Please run Step 3 — opportunity harvesting. Using the Compressed Client Dossier and the AI Readiness Snapshot, identify and score the top 5–7 AI opportunities for Meridian Talent Partners. Produce the full Scored Opportunity Map."

What to expect: An Executive Opportunity Summary paragraph, then 5–7 "opportunity cards" each with a description, pain point link, expected impact, feasibility notes, scores, classification, and a pilot suggestion. Ends with a Portfolio View grouping all opportunities into Quick Wins, Foundation Builders, and Big Bets.

After Claude finishes — trigger QG-3:

"Please run Quality Gate QG-3 on the Scored Opportunity Map."

4

Roadmap Composition

Produces the Recommended Action Sequence (1–2 pages, client-facing)

What this step does: Claude takes the scored opportunities and arranges them into a Now / Next / Later sequence (Months 1–3, 3–6, and 6–12). The sequencing follows strict rules — for example, data-heavy opportunities cannot go in the "Now" phase if the client's data maturity is "Early."

What to upload: Nothing new.

What to type:

"Please run Step 4 — roadmap composition. Using the Scored Opportunity Map and the AI Readiness Snapshot, sequence all opportunities into a Now / Next / Later action plan for Meridian Talent Partners. Produce the full Recommended Action Sequence."

What to expect: A Sequencing Rationale paragraph, then three phases (Now, Next, Later) each listing the relevant opportunities with a "why now/next/later" explanation grounded in maturity scores. Ends with a short "Bridge to Deeper Engagement" paragraph.

After Claude finishes — trigger QG-4:

"Please run Quality Gate QG-4 on the Recommended Action Sequence."

5

Blueprint Assembly

Produces the Final AI Value Blueprint (12–18 pages, client-facing)

What this step does: Claude takes all four prior outputs and compiles them into a single polished document — the final AI Value Blueprint that would be delivered to the client. This step produces a downloadable Word document (.docx file).

What to upload: Nothing new.

What to type:

"Please run Step 5 — Blueprint assembly. Using all four outputs from Steps 1 through 4, compile the final AI Value Blueprint for Meridian Talent Partners. Produce the complete client-facing deliverable as a .docx file."

What to expect: A complete 12–18 page document with 8 sections: Executive Summary, AI Readiness Snapshot, Key Findings, Opportunity Map, Action Sequence, Readiness Gaps, Recommended Next Steps, and Methodology Appendix. Claude will provide a download link for the .docx file.

After Claude finishes — trigger QG-FINAL (higher standard — requires Green, not just Amber):

"Please run QG-FINAL on the assembled Blueprint. Remember this gate requires Green (90%+) to pass — please check all 8 sections are present, all maturity scores match Step 2 exactly, all opportunity scores match Step 3 exactly, and no internal tags or methodology references are visible in the client-facing text."

Section 4: The Quality Gates — What They Mean and What to Do

After each step, Claude runs a quality check called a Quality Gate. This check scores the output and tells you whether it is good enough to move forward. Understanding what each result means — and what to do about it — is one of the most important parts of testing.

The four quality bands

GREEN 90%+	Output is strong. Move to the next step. For QG-1 through QG-4 this is ideal but not required — Amber is also acceptable. For QG-FINAL (Step 5), Green is the minimum required to deliver to a client.
AMBER 76%–89%	Output is acceptable with some issues. Claude will list the specific flags. Review each flag, decide if it needs fixing or is acceptable, then proceed to the next step. For QG-FINAL, Amber is not enough — you must remediate and re-run.
BLUE 60%–75%	Output has significant problems. Do not move forward yet. Ask Claude to fix the issues (see Section 6 for the exact words to use), then re-run the quality gate on the revised output.
RED Below 60%	Output is rejected. This is unusual with good input materials. If this happens, flag it as a test finding and see Section 6 for next steps.

What each quality gate is checking

Gate	After Step	Minimum	Key Things Being Checked
QG-1	Step 1 (Intake)	Amber (76%+)	Are all 8 dossier sections present? Are pain points backed by document evidence — not just form statements? Does every hypothesis have at least 2 evidence links?
QG-2	Step 2 (Maturity)	Amber (76%+)	Does every dimension score have at least one evidence citation from the dossier? Is conservative bias applied (ambiguous evidence → lower score)? Are key constraints named specifically?
QG-3	Step 3 (Opportunities)	Amber (76%+)	Does every opportunity trace to a specific named pain point? Are readiness adjustments applied and noted? Is the portfolio balanced (at least one Quick Win, one Foundation Builder, one Big Bet)?
QG-4	Step 4 (Roadmap)	Amber (76%+)	Are maturity gates respected (e.g., data-heavy opportunities not in Now if Data = Early)? Does Now include at least one Quick Win? Is every "why" grounded in a specific score or gap?
QG-FINAL	Step 5 (Assembly)	GREEN (90%+)	Is client name "Meridian Talent Partners" consistent throughout? Do all maturity scores match Step 2 exactly? Do all opportunity scores match Step 3 exactly? Are there zero internal tags visible in the client-facing text? Are all 8 sections present?

EXPECTED SCORES

For this test kit you should expect Green or Amber at QG-1 through QG-4, and Green at QG-FINAL. If you consistently see Blue or Red that is a significant finding worth noting — something in the pipeline needs investigation. Record all gate scores in Section 8.

Section 5: What Good Output Looks Like — Your Checklist

Use these checklists to review each step's output before triggering the quality gate. You do not need to understand every technical detail — these checks are things any consultant can verify with a careful read-through.

After Step 1 (Intake Dossier):

Step 1 Review Checklist
✓ The output mentions "Meridian Talent Partners" by name and describes them correctly (recruitment firm, Sofia, Bulgaria)
✓ The dossier has 8 lettered sections: A (Executive Summary), B (Key Data Points), C (Pain Points), D (Hypotheses), E (Org & Process), F (Document Index), G (Open Questions), H (Reviewer Checklist)
✓ Revenue of €4,823,400 appears in Section B and is tagged [Document-Backed]
✓ Time-to-fill of 38 days appears in Section B or C
✓ The CV formatting pain point (1.5–2 hours per candidate) appears in Section C
✓ The candidate database problem (47,000 records, ~35% stale) appears in Section C
✓ The GDPR constraint is mentioned in Section C or G
✓ Tags like [Document-Backed], [Form-Styled], [Inferred], [Assumption] appear throughout the text
✓ A "Quality Self-Assessment" block appears at the very end with scores for each of the 5 dimensions

After Step 2 (Maturity Snapshot):

Step 2 Review Checklist
✓ A table shows all 6 dimensions: Strategy, Data, Technology, People, Processes, Governance
✓ Each dimension is scored Early, Developing, or Established — not a number
✓ Data should be scored "Early" — evidence: 35% stale records, no data governance, GDPR gaps
✓ Governance should be scored "Early" — evidence: no AI ethics policy, GDPR consent records incomplete
✓ Technology should be scored "Developing" — evidence: Vincere migration underway, no system integrations
✓ People should be scored "Developing" at most — evidence: CEO open but senior Partners resistant
✓ Every score has a rationale paragraph with at least one citation to the dossier
✓ The "Key Constraints" section names at least 3 specific constraints (not generic statements)
✓ A "Quality Self-Assessment" block appears at the very end

After Step 3 (Opportunity Map):

Step 3 Review Checklist
✓ There are 5–7 opportunity cards — not more, not fewer

- ✓ Each opportunity card mentions a specific Meridian pain point — not a generic AI idea
- ✓ Look for these specific opportunities (or close equivalents): AI-assisted candidate sourcing, CV formatting automation, automated client status reporting, candidate database revival
- ✓ Opportunities are NOT vague generic phrases like "AI-Powered Analytics" with no connection to Meridian's actual situation
- ✓ Each card shows three scores: Impact (1–5), Feasibility (1–5), Alignment (1–5)
- ✓ Some feasibility scores should be reduced with a note like "Feasibility reduced due to Early Data maturity"
- ✓ GDPR appears as a note on any opportunity that involves candidate personal data
- ✓ The Portfolio View groups all opportunities into Quick Wins, Foundation Builders, and Big Bets
- ✓ A "Quality Self-Assessment" block appears at the very end

After Step 4 (Roadmap):

Step 4 Review Checklist

- ✓ Three phases appear: Now (Months 1–3), Next (Months 3–6), Later (Months 6–12)
- ✓ The Now phase includes at least one Quick Win — likely CV formatting automation or ChatGPT formalisation
- ✓ The Now phase does NOT include major database or AI sourcing work (Data is "Early" — this must wait)
- ✓ Any opportunity placed in Next or Later has a clear explanation of WHY it cannot start now
- ✓ That "why" explanation references a specific maturity gap, e.g., "Data maturity is Early, so this must follow the database governance work"
- ✓ A brief "Bridge to Deeper Engagement" paragraph appears at the end (consulting tone — not a sales pitch)
- ✓ A "Quality Self-Assessment" block appears at the very end

After Step 5 (Assembled Blueprint):

Step 5 Review Checklist

- ✓ The document is 12–18 pages long — count the sections in Claude's output or check the downloaded .docx file
- ✓ All 8 sections are present: Executive Summary, AI Readiness Snapshot, Key Findings, Opportunity Map, Action Sequence, Readiness Gaps, Recommended Next Steps, Appendix
- ✓ The name "Meridian Talent Partners" is used consistently throughout — not "Meridian" alone or "Meridian OOD" alone
- ✓ The maturity scores (Early/Developing/Established per dimension) exactly match what Step 2 produced
- ✓ The opportunity scores (Impact, Feasibility, Alignment numbers) exactly match what Step 3 produced
- ✓ The phase assignments (Now/Next/Later) exactly match what Step 4 produced
- ✓ NO internal tags are visible in the client text: no [Document-Backed], [Form-Stated], [Inferred], [Assumption]
- ✓ The tone throughout is professional and consultative — not overly negative about the client's situation

✓ The Executive Summary (page 1) is readable standalone — a CEO could read just that page and understand the core message

Section 6: What to Do When Something Goes Wrong

Problems during testing are expected and useful — they are findings, not failures. This section tells you how to handle the most common issues without needing technical help.

Problem: Quality Gate comes back Blue or Red

When this happens: Claude scores the output below 76% (Blue) or below 60% (Red).

What to do: Ask Claude to fix the specific issues it flagged. Type:

"The quality gate flagged the following issues: [paste the list Claude gave you]. Please revise the output to address each of these, then re-run the quality gate."

Note for your log: Most Blue results resolve in one revision. Red results are rare — if you see Red, note it as a significant finding and try once more before escalating.

Problem: A score or fact is wrong in a later step

When this happens: For example, Step 4 says a maturity score is "Established" but Step 2 scored it as "Developing".

What to do: Point the error out directly. Type:

"In the roadmap you described [dimension] as [wrong level], but the AI Readiness Snapshot from Step 2 scored it as [correct level]. Please correct this inconsistency."

Note for your log: Score drift between steps is one of the key things being tested. Always record it in your results log even after Claude fixes it.

Problem: The opportunities in Step 3 are generic

When this happens: An opportunity card says something like "AI-Powered Analytics Platform" with no link to a specific Meridian pain point.

What to do: Ask Claude to ground each opportunity in a specific pain point. Type:

"Opportunity #[X] does not connect to a specific Meridian pain point. Please revise it to trace directly to one of the pain points identified in the dossier, with a citation. If this opportunity cannot be grounded in the client's actual situation, remove it and replace it with one that can."

Note for your log: Generic opportunities are the most important quality failure to catch. Note each one in your results log.

Problem: Internal tags appear in the final Blueprint

When this happens: The assembled document (Step 5) contains phrases like [Document-Backed] or [Assumption] visible in the text.

What to do: Ask Claude to remove them. Type:

"The assembled Blueprint contains internal confidence tags that should not be visible to the client: [list the ones you found]. Please remove all such tags from the client-facing text and re-run QG-FINAL."

Note for your log: Any internal tags visible in the final document mean QG-FINAL cannot pass. This is an automatic failure condition.

Problem: Claude seems to have forgotten earlier outputs

When this happens: Claude starts making up new information or seems unaware of what it produced in earlier steps.

What to do: Remind Claude explicitly. Type:

"Please refer back to the outputs from Steps 1 through 4 that you produced earlier in this conversation. Do not introduce any new information — use only what was already established in the dossier, readiness snapshot, opportunity map, and action sequence."

Note for your log: If this persists despite the reminder, note it as a finding — it may indicate a context window limitation that needs to be addressed in the system design.

Section 7: Four Test Scenarios to Try

There are four different ways to run this test kit, each designed to test a different aspect of the system. Scenario A is the standard test — start there. The others are optional but will reveal different things about how the pipeline performs.

Scenario A

Full Pipeline — Standard Test

When to run it:	Start here. This is the primary test.
Files to upload:	Upload all 9 files (questionnaire answers + all 8 PDFs).
What to expect:	High-quality output throughout. Quality gates should mostly be Green or Amber. The final Blueprint should download as a clean, well-structured document. This establishes the baseline.
What to look for:	Do all 5 steps produce coherent, Meridian-specific output? Do the quality gates correctly identify any issues? Is the final document clean and professional?

Scenario B

Reduced Materials — Stress Test

When to run it:	After Scenario A, in a new conversation.
Files to upload:	Upload questionnaire answers + files 01–06 only. Do NOT upload file 07 (strategic plan) or file 08 (AI initiatives history).
What to expect:	The pre-pipeline screen should still say Sufficient (all 4 required files present). Some quality gate scores may drop — particularly on strategic alignment scoring in Step 3. You may see [Insufficient Evidence] flags where the pipeline cannot find strategic context.
What to look for:	Does the pipeline degrade gracefully? Does it correctly flag gaps rather than making things up? Do the confidence scores drop appropriately?

Scenario C

Form Only — Thin Materials Test

When to run it:	After Scenario A, in a new conversation.
Files to upload:	Upload ONLY the questionnaire answers file (00). Do not upload any PDFs.
What to expect:	The pre-pipeline screen should say Marginal or Insufficient. If you proceed anyway, Step 1 should produce a dossier where almost all claims are tagged [Form-Stated] rather than [Document-Backed]. Quality gates should score lower.
What to look for:	Does the pre-pipeline screen correctly flag that required documents are missing? Does Step 1 honestly reflect the lower evidence quality? Does it avoid making things up to fill gaps?

Scenario D

GDPR Sensitivity Check

When to run it:	Can be done as part of Scenario A — just pay extra attention during Steps 3 and 4.
Files to upload:	Same as Scenario A (all files).
What to expect:	The pipeline should NOT recommend "build an AI matching system using your 47,000-record candidate database" in the Now phase. It should flag GDPR as a prerequisite — the database cannot be used for AI applications until a data governance and GDPR consent audit is completed first.
What to look for:	Does the pipeline correctly identify GDPR as a feasibility constraint in the maturity scoring (Governance = Early), the opportunity scoring (reduced feasibility for database-dependent ideas), and the roadmap sequencing (database work in Next, not Now)?

Section 8: Recording Your Results

Please record your findings as you test. This information is valuable — it tells us where the system is working well and where it needs improvement. You do not need to write lengthy reports. The tables below are all you need to fill in.

Quality Gate Score Log

Fill this in for each pipeline run. Record the score and band Claude reports for each gate.

Gate	After Step	Score (%)	Band (Green / Amber / Blue / Red)	Notes / Issues Flagged
Pre-Pipeline	Before Step 1	_____ %	_____	
QG-1	Step 1 (Intake)	_____ %	_____	
QG-2	Step 2 (Maturity)	_____ %	_____	
QG-3	Step 3 (Opportunities)	_____ %	_____	
QG-4	Step 4 (Roadmap)	_____ %	_____	
QG-FINAL	Step 5 (Assembly)	_____ %	_____	

Key Facts Verification Log

Check whether these specific numbers from the test data appear correctly in the pipeline output. Tick Yes if the number appears and is correct, No if it is wrong or missing, N/A if it was not relevant at that step.

Fact to Check	Expected Value	In Step 1?	In Step 2?	In Step 3?	In Step 5?
Total FY2025 revenue	€4,823,400	■ Yes ■ No	■ N/A	■ N/A	■ Yes ■ No
Average time-to-fill	38 days	■ Yes ■ No	■ N/A	■ Yes ■ No	■ Yes ■ No
Sourcing hours per mandate	6–8 hours	■ Yes ■ No	■ N/A	■ Yes ■ No	■ Yes ■ No
Candidate database records	~47,000 records	■ Yes ■ No	■ N/A	■ Yes ■ No	■ Yes ■ No
CV formatting time	1.5–2 hrs per cand.	■ Yes ■ No	■ N/A	■ Yes ■ No	■ Yes ■ No
Client satisfaction — comms	3.1 / 5.0	■ Yes ■ No	■ N/A	■ Yes ■ No	■ Yes ■ No
Data dimension maturity score	Early	■ N/A	■ Yes ■ No	■ Yes ■ No	■ Yes ■ No

Fact to Check	Expected Value	In Step 1?	In Step 2?	In Step 3?	In Step 5?
Governance dimension maturity score	Early	■ N/A	■ Yes ■ No	■ Yes ■ No	■ Yes ■ No
GDPR mentioned as constraint	Yes — must appear	■ Yes ■ No	■ Yes ■ No	■ Yes ■ No	■ Yes ■ No
ChatGPT flagged as compliance risk	Yes — must appear	■ Yes ■ No	■ N/A	■ Yes ■ No	■ Yes ■ No

Issues Found Log

Use this to record any problems, unexpected behaviour, or findings during your test.

Step	Description of Issue	Severity (High / Med / Low)	Resolved?
			■ Yes ■ No
			■ Yes ■ No
			■ Yes ■ No
			■ Yes ■ No
			■ Yes ■ No
			■ Yes ■ No
			■ Yes ■ No
			■ Yes ■ No

Overall Summary

Scenario tested (A / B / C / D):	
Date and time of test:	
Total time for full pipeline run:	
Gates that required remediation:	

Was the final Blueprint downloaded successfully?	☐ Yes ☐ No
Did the final Blueprint look client-ready?	☐ Yes ☐ No ☐ Partially
Overall verdict:	☐ Working well ☐ Issues to investigate ☐ Major problems
Most important finding from this test:	

 **SEND
YOUR
RESULTS**

When you have finished testing, please send your completed results log to the project lead. Even a single finding — one wrong number, one generic opportunity, one quality gate score that surprised you — is useful. There are no wrong answers. Your observations are exactly what this testing process is for.